

```

import math

from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
import numpy as np

np.random.seed(1337)

X_SIZE = 5
L = 100

BETA_LIST = (0.5, 1, 1, 1, 1, 1)

N_LIST = (
    [number for number in range(50, 100, 10)] +
    [number for number in range(100, 1100, 100)]
)

def generate_df(n: int, x_size: int = X_SIZE) -> np.array:
    return np.array([
        [np.random.normal(0, 1) for _ in range(x_size)]
        for _ in range(n)
    ])

def probability(x: list[float], b_list: list[float]) -> float:
    b_list = list(b_list)

    b0 = b_list.pop(0)

    exp_content = b0
    for i in range(len(x)):
        exp_content += b_list[i]*x[i]

    return 1 / (1 + math.exp(-exp_content))

def mse(beta_head: list[float], beta: list[float]) -> float:
    beta_head = np.array(beta_head)
    beta = np.array(beta)

    return np.mean((beta_head - beta)**2)

```

1. Fit logistic model and calculate the estimators of the coefficients  $\beta = (\beta_1, \beta_2, \beta_3, \beta_4, \beta_5)$ . Repeat the experiment  $L = 100$  times and compute the MSE

```

results: dict[int, list[float]] = {}

for n in N_LIST:

```

```

results[n] = []
for _ in range(L):
    df = generate_df(n)

    p_list = [
        probability(x, BETA_LIST)
        for x in df
    ]

    y_list = [
        np.random.binomial(1, p)
        for p in p_list
    ]

    lf = LogisticRegression(penalty="l2", C=1000)
    lf.fit(df, y_list)

    beta_head = lf.coef_

    mse_value = mse(beta_head, BETA_LIST[1:])

    results[n].append(mse_value)

round_value = 3

for key in results:
    min_v = round(min(results[key]), round_value)
    avg_v = round(np.mean(results[key]), round_value)
    max_v = round(max(results[key]), round_value)

    print(f"{key:7} {min_v:7} {avg_v:7} {max_v:7}")

```

50	0.014	5.128	233.857
60	0.019	0.401	6.588
70	0.012	0.352	5.122
80	0.003	0.284	3.694
90	0.01	0.206	2.006
100	0.013	0.183	1.58
200	0.007	0.074	0.589
300	0.005	0.046	0.175
400	0.004	0.026	0.083
500	0.003	0.025	0.118
600	0.002	0.017	0.072
700	0.002	0.015	0.084
800	0.001	0.013	0.064
900	0.001	0.011	0.041
1000	0.003	0.009	0.026

```

def draw(
    results: dict[int, list[float]],
    draw_first: bool = True,

```

```

        max_value: float = None
    ) -> None:

    if max_value is None:
        max_value = float("inf")

    dict_to_plot = {
        key: [value for value in values if value <= max_value]
        for key, values in results.items()
    }

    if not draw_first:
        min_key = min(list(results.keys()))
        del dict_to_plot[min_key]

    labels, data = zip(*dict_to_plot.items())

    plt.figure(figsize=(15, 10))

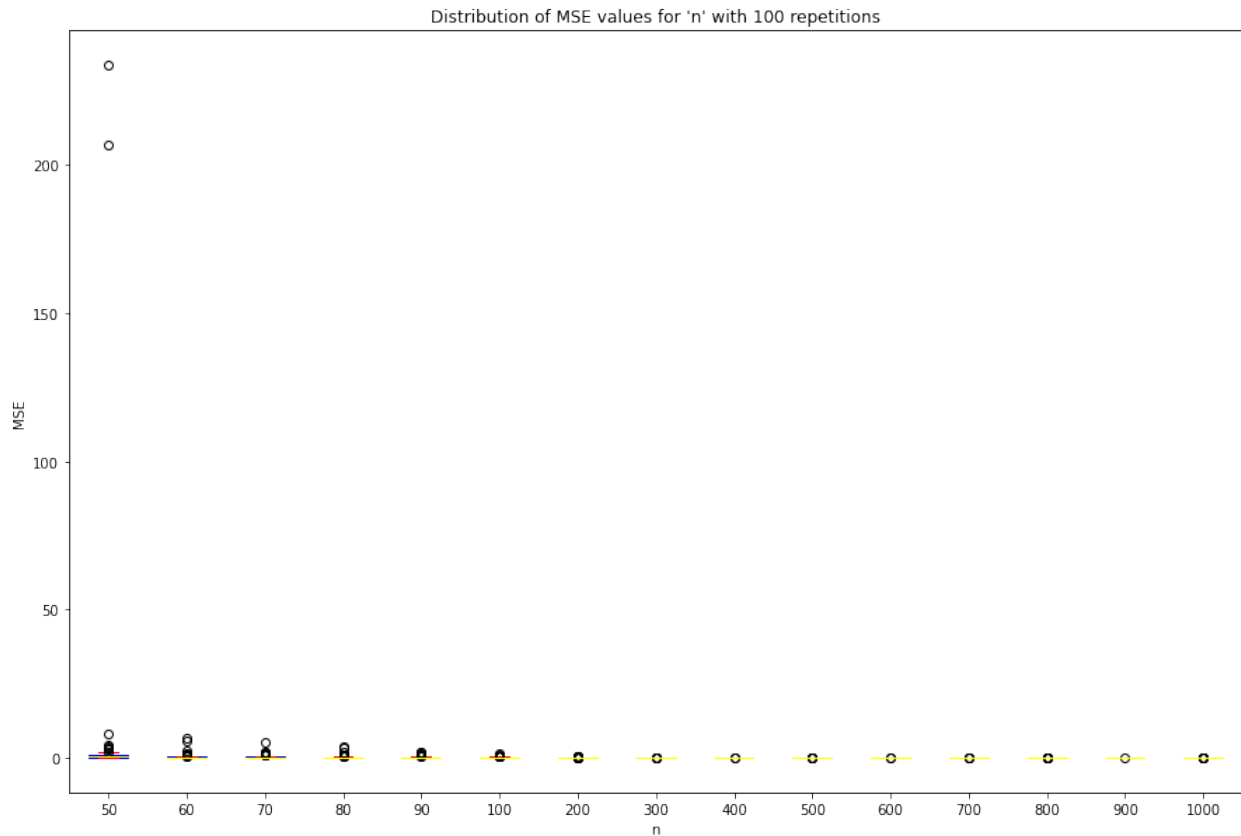
    plt.boxplot(
        data, patch_artist=True, labels=labels,
        boxprops={"facecolor": "lightblue", "color": "blue"},
        whiskerprops={"color": "green"},
        capprops={"color": "red"},
        medianprops={"color": "yellow"}
    )

    plt.title("Distribution of MSE values for 'n' with 100
repetitions")
    plt.xlabel("n")
    plt.ylabel("MSE")

    plt.show()

draw(results)


```



As we can see, for small values (as 50) the MSE could be a really high.

Let's remove it for greater readability of the chart.


```
draw(results, draw_first=False)
```



Now it is better. However, there are a few outsiders in the data for smaller values.

Let's set the limit at Y axis, for example 1.

```
draw(results, max_value=1)
```



Now it is readable.

As we can see, the MSE is lower and lower with biggest 'n' number.

That means, that more samples means lower error. If we want to have a really good model, we need to have a lot of data for fitting

2. Using the same datasets, train the model based only on 3 variables:  $x_1, x_2, x_3$  and draw the analogous curve showing how MSE for  $\beta = (\beta_1, \beta_2, \beta_3)$  depends on  $n$ .

```
BETA_LIST = (0.5, 1, 1, 1)
new_results: dict[int, list[float]] = {}

for n in N_LIST:
    new_results[n] = []
    for _ in range(L):
        df = generate_df(n, 3)

        p_list = [
            probability(x, BETA_LIST)
            for x in df
        ]
```

```

y_list = [
    np.random.binomial(1, p)
    for p in p_list
]

lf = LogisticRegression(penalty="l2", C=1000)
lf.fit(df, y_list)

beta_head = lf.coef_

mse_value = mse(beta_head, BETA_LIST[1:])

new_results[n].append(mse_value)

round_value = 3


for key in new_results:
    min_v = round(min(new_results[key]), round_value)
    avg_v = round(np.mean(new_results[key]), round_value)
    max_v = round(max(new_results[key]), round_value)

    print(f"{key:7} {min_v:7} {avg_v:7} {max_v:7}")

```

50	0.008	0.367	4.527
60	0.002	0.238	1.378
70	0.003	0.226	3.637
80	0.004	0.175	1.19
90	0.002	0.11	0.708
100	0.005	0.109	0.567
200	0.004	0.049	0.275
300	0.001	0.027	0.17
400	0.001	0.021	0.076
500	0.001	0.019	0.086
600	0.0	0.017	0.112
700	0.001	0.012	0.076
800	0.001	0.013	0.053
900	0.0	0.01	0.054
1000	0.0	0.009	0.032

```
draw(new_results)
```




Now we have only three X variables.

There are still some outsiders, lets set Y limit to 1.

```
draw(new_results, max_value=1)
```





The plot looks really similar to the last one with five X variables.

Lets compare them correctly.


```
plt.figure(figsize=(20, 10))

averages = {label: sum(values) / len(values) for label, values in
results.items() if label != 50}
plt.plot(list(averages.keys()), list(averages.values()), marker="o",
label="5 values")

averages = {label: sum(values) / len(values) for label, values in
new_results.items() if label != 50}
plt.plot(list(averages.keys()), list(averages.values()), marker="*",
label="3 values")

plt.xticks([50] + list(range(100, 1100, 100)))
plt.title("Average MSE value depending on the number of 'beta' and
'n'")
plt.xlabel("n")
plt.ylabel("AVG MSE")
plt.legend()

plt.show()
```



You can see that at first the model learns easily for fewer variables - which sounds logical.

But with more samples loaded into the model, the number of variables becomes less important. The model trains just as well.